

## Geoint Accelerator

- Platform to start developing with GPUs
- GPU access, GPU-accelerated apps and libraries
- Register to learn more: <http://goo.gl/Eui6k6>

## Webinar Feedback

Submit your feedback for a chance to win Tesla K20 GPU\*

[https://www.surveymonkey.com/s/OpenCV\\_July30](https://www.surveymonkey.com/s/OpenCV_July30)


\* Offer is valid until August 13<sup>th</sup> 2013



More questions on OpenCV and GPUs

- Stay tuned with NVIDIA webinars:  
[http://www.nvidia.com/object/cuda\\_signup\\_alerts.html](http://www.nvidia.com/object/cuda_signup_alerts.html)
- Refer to OpenCV Yahoo! Groups
- OpenCV.org: <http://answers.opencv.org/questions/>

# USING NVIDIA GPUS WITH OPENCV FOR ACCELERATED HIGH PERFORMANCE COMPUTER VISION



Anatoly Baksheev, Itseez Inc.



# Outline

- Getting and building OpenCV with CUDA
- GPU module API
- Overlapping operations
- Using GPU module with your CUDA code
- Questions & Answers



# What is OpenCV

- Popular Computer Vision library
  - 6M downloads
  - BSD license
  - 1000s CV functions for solving various problems
  - Various optimizations
    - **OpenCV GPU** module contains **CUDA** acceleration
    - Intel TBB, IPP, SSE, ARM NEON & GLSL (Tegra)



# Getting OpenCV

- Source Forge
  - <http://sourceforge.net/projects/opencvlibrary/>
    - Source package for each release (**recommended**)
    - OpenCV GPU Demo Pack binaries
    - (Precompiled OpenCV binaries with CUDA support)
- Git repository (modified daily)
  - `git clone git://github.com/Itseez/opencv.git` (GitHub mirror)
  - `git checkout -b 2.4 origin/2.4`
- Two branches
  - **2.4** with binary compatibility between releases (tagged as 2.4.x)
  - **master** (coming OpenCV 3.0 with algorithms API changes, [separate webinar](#))



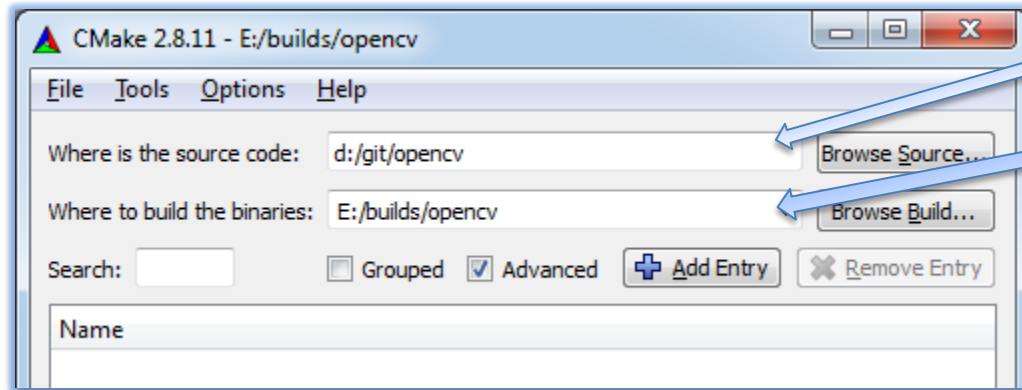
# Building OpenCV with GPU support

- Prerequisites
  - OpenCV sources
  - CMake 2.8.8 (<http://www.cmake.org/>)
    - CMake-GUI
  - NVIDIA Display Driver
  - NVIDIA GPU Computing Toolkit (for CUDA)
    - <https://developer.nvidia.com/cuda-toolkit>
  - And your favorite IDE/compiler



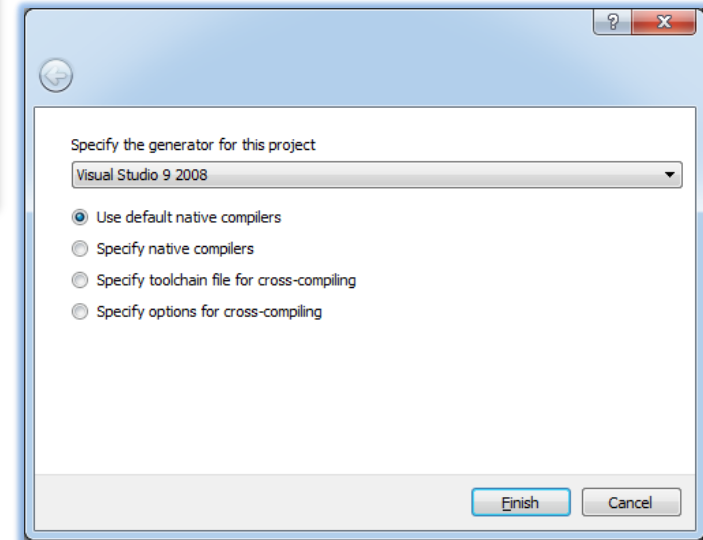

# Building OpenCV with GPU support

- Build steps (screenshots for Windows 7, Visual Studio)
  - Run CMake GUI and set **source** and **build** directories, press **Configure** and select you compiler to generate project for.



The base folder of OpenCV source code

Where to put compiled OpenCV library



# Building OpenCV with GPU support

- Build steps
  - Run CMake GUI and set **source** and **build** directories, press **Configure** and select you compiler to generate project for.
  - Enable **WITH\_CUDA** flag and ensure that CUDA Toolkit is detected correctly by checking all variables with 'CUDA\_' prefix.

Name	Value
CUDA_TOOLKIT_INCLUDE	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0/include
CUDA_TOOLKIT_ROOT_DIR	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0
CUDA_VERBOSE_BUILD	<input type="checkbox"/>
CUDA_VERSION	5.0
CUDA_cublas_LIBRARY	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0/lib/Win32/cublas.lib
CUDA_cufft_LIBRARY	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0/lib/Win32/cufft.lib
CUDA_cupti_LIBRARY	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0/extras/CUPTI/libWin32/cupti.lib
CUDA_curand_LIBRARY	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0/lib/Win32/curand.lib
CUDA_cuspars_LIBRARY	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0/lib/Win32/cuspars.lib
CUDA_npp_LIBRARY	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0/lib/Win32/npp.lib
CUDA_nvccuenc_LIBRARY	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0/lib/Win32/nvccuenc.lib
CUDA_nvccuid_LIBRARY	C:/Program Files/NVIDIA GPU Computing Toolkit/CUDA/v5.0/lib/Win32/nvccuid.lib
WITH_CUDA	<input checked="" type="checkbox"/>

Press Configure to update and display new values in red, then press Generate to generate selected build files.





# Building OpenCV with GPU support

- Build steps
  - Run CMake GUI and set **source** and **build** directories, press **Configure** and select you compiler to generate project for.
  - Enable **WITH\_CUDA** flag and ensure that CUDA Toolkit is detected correctly by checking all variables with 'CUDA\_' prefix.
  - Press **Configure** and **Generate** to generate a project
    - On Windows, open the Visual Studio solution and click on “Build Solution”.
    - On Linux, run “make” from the build folder.



# Installing OpenCV (optional)

- Running OpenCV install scripts is a way to put all headers, libs and binaries to one place for easier use and deployment
  - Set **CMAKE\_INSTALL\_PREFIX** variable
  - Build **INSTALL** target

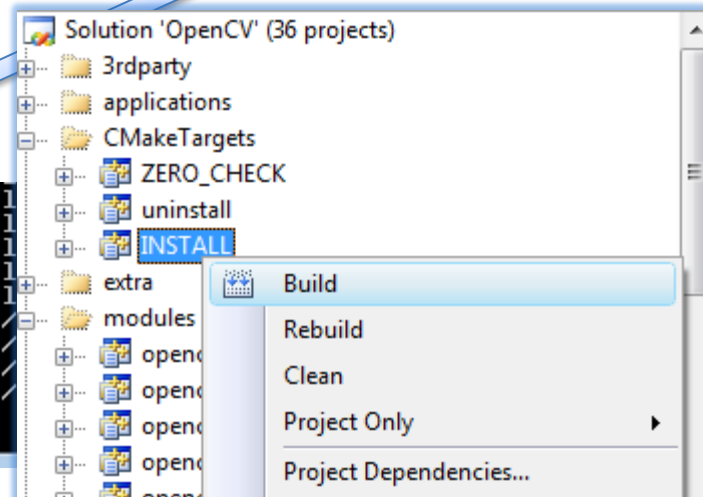
Where is the source code:

Where to build the binaries:


Search:

Name	Value
CMAKE_INSTALL_PREFIX	E:/builds/opencv/install
CMAKE_SKIP_INSTALL_RPATH	<input type="checkbox"/>
CUDA_HOST_COMPILER	\$(VCInstallDir)bin
INSTALL_C_EXAMPLES	<input type="checkbox"/>

```
-- Up-to-date: E:/opencv24_qtcr/install/incl
-- Up-to-date: E:/opencv24_qtcr/install/incl
-- Up-to-date: E:/opencv24_qtcr/install/incl
-- Up-to-date: E:/opencv24_qtcr/install/bin/
-- Up-to-date: E:/opencv24_qtcr/install/bin/
-- Up-to-date: E:/opencv24_qtcr/install/bin/
E:\opencv24_qtcr>make install
```



Binaries



Include folder for all (!) modules

all libs (!)



# Adding OpenCV to your project via CMake

- Create a 'main.cpp' file and **CMakeLists.txt** script as below

```
1  cmake_minimum_required(VERSION 2.8.8)
2
3  project(sample_project C CXX)
4
5  find_package(OpenCV REQUIRED)
6  include_directories(${OpenCV_INCLUDE_DIRS})
7
8  file(GLOB srcs ./*.cpp ./*.h*)
9
10 add_executable(sample_app ${srcs})
11 target_link_libraries(sample_app ${OpenCV_LIBS})
```

- Use CMake to generate and build the sample. Set **OpenCV\_DIR** variable for the project to **build** or **install** OpenCV folder

Name	Value
OpenCV_3RDPARTY_LIB_DIR_DBG	E:/builds/opencv/install/share/OpenCV/3rdparty/lib
OpenCV_3RDPARTY_LIB_DIR_OPT	E:/builds/opencv/install/share/OpenCV/3rdparty/lib
OpenCV_CONFIG_PATH	E:/builds/opencv/install
OpenCV_DIR	E:/builds/opencv/install
OpenCV_LIB_DIR_DBG	E:/builds/opencv/lib
OpenCV_LIB_DIR_OPT	E:/builds/opencv/installlib

Set this. Others  
are detected automatically



# Outline

- Getting and building OpenCV with CUDA
- GPU module API
- Overlapping operations
- Using GPU module with your CUDA code
- Questions & Answers



# GPU module design considerations

- Key ideas
  - Explicit control of data transfers between CPU and GPU
  - Minimization of the data transfers
  - Completeness
    - Port everything even functions with little speed-up
- Solution
  - Container for GPU memory with upload/download functionality
  - GPU module function take the container as input/output parameters



# GpuMat - container for GPU memory

- Class GpuMat – for storing 2D (**pitched**) data on GPU
  - Interface similar to **cv::Mat()**, supports reference counting
  - Its data is not continuous, extra **padding** in the end of each row
  - It contains:
    - **data** - Pointer data beginning in GPU memory
    - **step** - distance **in bytes** is between two consecutive rows
    - **cols, rows** – fields that contain image size
    - Other fields for internal use only

Can use with  
another  
CUDA libs

```
GpuMat cloud = ... (get 3-channel float data from a file and copy to GPU)

// compute address of Point at (x, y)
cv::Point3f* pixel_ptr = (cv::Point3f*)(cloud.data + cloud.step * y) + x;

// the same, but simplified
cv::Point3f* pixel_ptr = cloud.ptr<cv::Point3f>(y) + x;
```



# Operations with GpuMat

- Allocations (similar to cv::Mat)

```
void GpuMat::GpuMat(const cv::Size& size, int type);  
void GpuMat::create(const cv::Size& size, int type);  
  
//Type examples:  
    CV_8U - grayscale,  
    CV_8UC3 - BGR,  
    CV_32FC3 - 3D point  
    CV_16U - depth  
  
GpuMat image1(Size(1902,1080), CV_8U);  
  
GpuMat image2;  
image2.create(Size(1, 1000), CV_32FC3);
```

- Creating GpuMat **header** for user allocated data

```
void GpuMat::GpuMat(const cv::Size& size, int type, void* data, size_t step);  
  
thrust::device_vector<float> vector(10000);  
GpuMat header(Size(vector.size(), 1), CV_32F, &vector[0], vector.size() * sizeof(float));  
cv::gpu::threshold(header, header, value, 1000, THRESH_BINARY);
```



# Operations with GpuMat

- Copying between CPU-GPU and GPU-GPU

```
void GpuMat::GpuMat(const cv::Mat& host_data);  
void GpuMat::upload(const cv::Mat& host_data);  
  
void GpuMat::download(cv::Mat& host_data);  
void GpuMat::copyTo(cv::gpu::GpuMat& other);  
  
//Examples  
Mat host_image = cv::imload("file.png"); //load image from file  
  
GpuMat device_image1;  
device_image1.upload(host_image);           //allocate memory and upload to GPU  
  
GpuMat device_image2;  
device_image1.copyTo(device_image2);        //allocate memory and GPU-GPU copy  
  
device_image2.download(host_image);         //download data
```





# Template matching

```
#include <opencv2/opencv.hpp>
#include <opencv2/gpu/gpu.hpp>

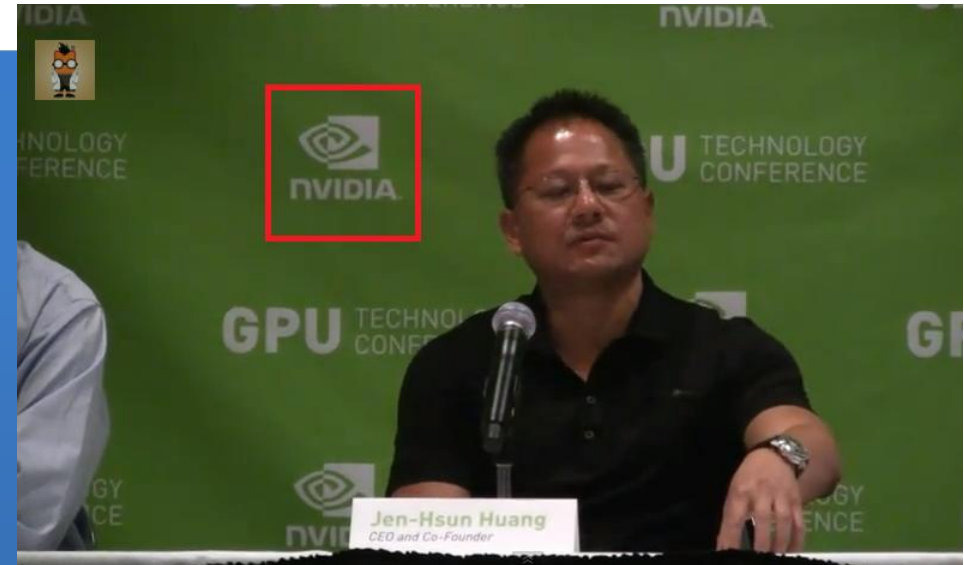
void process() {
    cv::gpu::setDevice(0);
    cv::VideoCapture capture("video.avi");
    cv::Mat templ_h = cv::imread("nvlogo.png");
    cv::gpu::GpuMat templ_d(templ_h); // upload
    cv::gpu::GpuMat image_d, result;
    cv::Mat image_h;

    for(;;) {
        capture >> image_h;
        if(image_h.empty())
            break;

        image_d.upload(image_h);
        cv::gpu::matchTemplate(image_d, templ_d, result, cv::TM_CCORR);

        double max_value;
        cv::Point location;
        cv::gpu::minMaxLoc(result, 0, &max_value, 0, &location);
        if (max_value > threshold)
            drawRectangleAndShowImage(location, ...);
    }
}
```

Init CUDA



Template  
searched



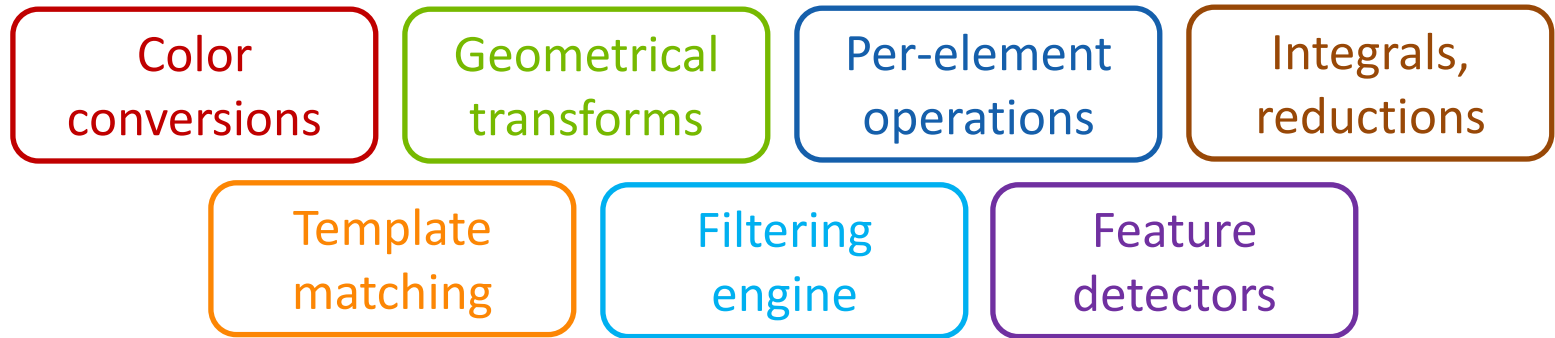
8-20x

\* GTX 680 vs. Core i5-760




# OpenCV GPU Module functionality

- Image processing building blocks:



- High-level algorithms:



# Outline

- Getting and building OpenCV with CUDA
- GPU module API
- **Overlapping operations**
- Using GPU module with your CUDA code
- Questions & Answers



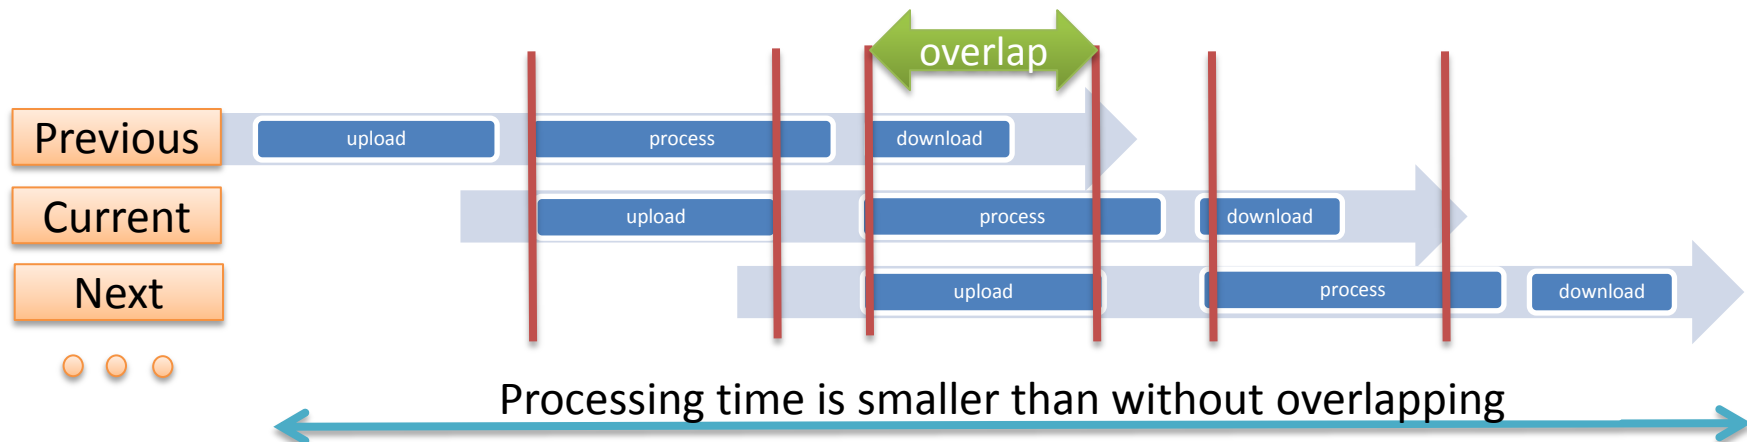
# Overlapping operations with CUDA

- Which operations can be run simultaneously?
  - Host and device code execution
  - Data transfer and device code execution
  - Several device code executions
    - Only if GPU is under occupied
    - Must be executed in different streams
    - (won't talk about this)



# Overlapping operations with CUDA

- Advantages
  - Allows to use whole computing power of your PC (CPU + GPU)
  - Allows to partly **hide data transfer overheads**
    - Result from previous frame is being downloaded
    - Current is being processed on GPU
    - Next frame is being uploaded



# Concurrency in OpenCV GPU

- Class CudaMem
  - Allocates page-locked CPU memory
    - Required for asynchronous data transfers,
      - **2x faster** transfer than regular memory, but limited by available RAM
  - Interface similar to Mat or GpuMat
  - Convertible to Mat, can pass to any CPU function

```
CudaMem page_locked(Size(1920, 1080), CV_32FC3); //allocate page locked memory
Mat header = page_locked;                        //no copy, just header

VideoCapture cap("video1080p.avi"); // open video stream
for(;;)
{
    cap >> header; //read video frame to the header which actually points to page-locked buffer
    // now can copy the page_locked to GPU in asynchronous fashion
    ...
}
```



# Concurrency in OpenCV GPU

- Class Stream (wrapper for cudaStream\_t)
  - Represents asynchronous queue of operations
  - Allows querying/waiting operations to complete
  - Enqueue asynchronous data transfers

```
bool Stream::queryIfComplete();  
void Stream::waitForCompletion();  
  
void Stream::enqueueDownload(const GpuMat& src, CudaMem& dst);  
void Stream::enqueueUpload(const CudaMem& src, GpuMat& dst);  
void Stream::enqueueCopy(const GpuMat& src, GpuMat& dst);  
  
typedef void (*StreamCallback)(Stream& stream, int status, void* userData);  
void Stream::enqueueHostCallback(StreamCallback callback, void* userData);
```



# Concurrency in OpenCV GPU

- Putting device operations into the queue

```
namespace cv { namespace gpu
{
    void a_opencv_gpu_function(..., Stream& stream = Stream::Null());
}}

// synchronous case
device_image.upload(host_image);           // blocks until upload is done
a_opencv_gpu_function(device_image, output); // blocks until operation finishes

// asynchronous case
Stream stream;

stream.enqueueUpload(host_image, device_image); // returns immediately
a_opencv_gpu_function(device_image, output, stream); // returns immediately
Stream.enqueueDownload(output, output_host); // returns immediately

// CPU resources are available.
// Let's compute on CPU here while GPU does own work.

stream.waitForCompletion();
// output_host is ready!
```





# Concurrency in OpenCV GPU

- Current limitation:
  - Unsafe to enqueue the same GPU operation multiple times

```
__constant__ int variable; // constant GPU memory reference
texture<uchar, 2> tex_reference;

cv::gpu::function1(int param, Stream& stream)
{
    cudaMemCopyToSymbol (variable, param); // variable = param
    call_gpu_kernel_that_uses_variable_in_async_fashion(stream)
}

//unsafe
function1(10, stream);
function1(20, stream);
```

- The limitation will be removed in OpenCV 3.x
  - (will re-implement without using constant memory although it may lead to slightly worse performance for small number of kernels)



# Outline

- Getting and building OpenCV with CUDA
- GPU module API
- Overlapping operations
- Using GPU module with your CUDA code
- Questions & Answers



# Writing your own CUDA code

- GpuMat (can't be passed to cu-file due to nvcc compiler issue, this will be fixed in OpenCV 3.0)
  - data, step, cols, rows – can just pass to your code
  - Convertible to PtrStep<T>, PtrStepSz<T> structures

```
// swap_rb.cpp
#include <opencv2/gpu/stream_accessor.hpp>

void swap_rb_caller(const PtrStepSz<uchar3>& src, PtrStep<uchar3> dst, cudaStream_t stream);

void swap_rb(const GpuMat& src, GpuMat& dst, Stream& stream = Stream::Null())
{
    CV_Assert(src.type() == CV_8UC3);
    dst.create(src.size(), src.type()); // create if not allocated yet
    cudaStream_t s = StreamAccessor::getStream(stream);
    swap_rb_caller(src, dst, s);
}
```



# Writing your own CUDA code

```
// swap_rb.cu
#include <opencv2/core/cuda_devptrs.hpp>

__global__ void swap_rb_kernel(const PtrStepSz<uchar3> src, PteStep<uchar3> dst)
{
    int x = threadIdx.x + blockIdx.x * blockDim.x;
    int y = threadIdx.y + blockIdx.y * blockDim.y;

    if (x < src.cols && y < src.rows)
    {
        uchar3 v = src(y, x); // Reads pixel in GPU memory. Valid! We are on GPU!
        dst(y, x) = make_uchar3(v.z, v.y, v.x);
    }
}

void swap_rb_caller(const PtrStepSz<uchar3>& src, PtrStep<uchar3> dst, cudaStream_t stream)
{
    dim3 block(32, 8);
    dim3 grid((src.cols + block.x - 1)/block.x, (src.rows + block.y - 1)/ block.y);

    swap_rb_kernel<<<grid, block, 0, stream>>>>(src, dst);

    if (stream == 0)
        cudaDeviceSynchronize();
}
```



# Writing your own CUDA code

- Device layer functions (the set will be extended in OpenCV 3.0)
  - opencv2/gpu/device/\*.hpp
  - useful utility code to call from your kernels

```
#define HIST_SIZE 396
#include <opencv2/gpu/device/block.hpp>
__global__ void histogramm_kernel_pass1(const PtrStepSz<float3> cloud, PtrStep<float> ouput) {
    __shared__ float smem[HIST_SIZE];

    Block::fill(smem, smem + HIST_SIZE, 0);
    __syncthreads()

    ... compute histogram using atomics on smem ...

    __syncthreads();
    float sum = Block::reduce_n(smem, HIST_SIZE, plus<float>());
    __syncthreads();

    float* ouput_row = output(blockIdx.y * BLOCK_SIZE + blockIdx.x);
    Block::transform(smem, smem + HIST_SIZE, output_row, divide_by(sum));
}
```



# Writing your own CUDA code

- Device layer functions (the set will be expanded)
  - opencv2/gpu/device/\*.hpp
  - useful utility code to call from your own code

```
#define HIST_SIZE 396
#include <opencv2/gpu/device/block.hpp>
__global__ void histogramm_kernel_pass1(const Ptr_StepSmem,
__shared__ float smem[HIST_SIZE];

    Block::fill(smem, smem + HIST_SIZE, 0);
    __syncthreads()

    ... compute histogram using atomics on smem ...

    __syncthreads();
    float sum = Block::reduce_n(smem, HIST_SIZE, plus<float>());
    __syncthreads();

    float* output_row = output(blockIdx.y * BLOCK_SIZE + blockIdx.x);
    Block::transform(smem, smem + HIST_SIZE, output_row, divide_by(sum));
}
```


```
template <typename T, class BinOp>
static __device__ __forceinline__
void reduce_n(T* data, unsigned int n, BinOp op)
{
    int ftid = threadIdx.z * blockDim.x * blockDim.y;
    int sft = blockDim.x * blockDim.y * blockDim.z;
    if (sft < n)
    {
        for (unsigned int i = sft + ftid; i < n; i += sft)
            data[ftid] = op(data[ftid], data[i]);
        __syncthreads();
        n = sft;
    }

    while (n > 1)
    {
        unsigned int half = n/2;
        if (ftid < half)
            data[ftid] = op(data[ftid], data[n - ftid - 1]);
        __syncthreads();
        n = n - half;
    }
}
```



# OpenCV resources

- Main site:
  - <http://opencv.org/>
- Online documentation (**with tutorials**)
  - <http://docs.opencv.org/>
- Question & Answers site
  - <http://answers.opencv.org/questions/>
- Books
  - <http://opencv.org/books.html>
- Bug tracker
  - <http://code.opencv.org/projects/opencv/issues>



“**Learning OpenCV**” 2<sup>nd</sup> ed.  
by G. Bradski, A. Kaehler  
with **GPU module** chapter  
coming soon!



# Questions?

[www.opencv.org](http://www.opencv.org)





# Upcoming GTC Express Webinars

**July 31** - NMath Premium: GPU-Accelerated Math Libraries for .NET

**August 7** - Accelerating High Performance Computing with GPUDirect RDMA

**August 13** - GPUs in the Film Visual Effects Pipeline

**August 14** - Beyond Real-time Video Surveillance Analytics with GPUs

**Register at [www.gputechconf.com/gtcexpress](http://www.gputechconf.com/gtcexpress)**

# GTC 2014 Call for Submissions

Looking for submissions in the fields of

- Science and research
- Professional graphics
- Mobile computing
- Automotive applications
- Game development
- Cloud computing



Submit at [www.gputechconf.com](http://www.gputechconf.com)

## Geoint Accelerator

- Platform to start developing with GPUs
- GPU access, GPU-accelerated apps and libraries
- Register to learn more: <http://goo.gl/Eui6k6>

## Webinar Feedback

Submit your feedback for a chance to win Tesla K20 GPU\*

[https://www.surveymonkey.com/s/OpenCV\\_July30](https://www.surveymonkey.com/s/OpenCV_July30)

\* Offer is valid until August 13<sup>th</sup> 2013



More questions on OpenCV and GPUs

- Stay tuned with NVIDIA webinars:  
[http://www.nvidia.com/object/cuda\\_signup\\_alerts.html](http://www.nvidia.com/object/cuda_signup_alerts.html)
- Refer to OpenCV Yahoo! Groups
- OpenCV.org: <http://answers.opencv.org/questions/>